

HPC LAB : 02
OpenMp Directives, Spmd, worksharing

PRN : 2019BTECS00055
Name : Aditya Rajendra Manapure

Problem Statement 01 . Spmd vs worksharing in OpenMp

WorkSharing : Using the OpenMP pragmas, most loops with no loop-carried dependencies can be threaded with one simple statement. This concept of OpenMP to parallelize loops, is called **worksharing**. A worksharing region has no barrier on entry; however, an implied barrier exists at the end of the worksharing region, unless a clause **nowait** is specified. If a **nowait** clause is present, an implementation may omit the barrier at the end of the worksharing region. In this case, threads that finish early may proceed straight to the instructions that follow the worksharing region without waiting for the other members of the team to finish the worksharing region, and without performing a flush operation.

SPMD : Computers in which each processing element is capable of executing a different program independent of the other processing elements are called **multiple instruction stream, multiple data stream MIMD** computers. A simple variant of this model, called the **single program multiple data SPMD** model, relies on multiple instances of the same program executing on different data.

Problem Statement 01 : OpenMP Clauses

For general attributes:

Clause	Description
if	Specifies whether a loop should be executed in parallel or in serial.
num_threads	Sets the number of threads in a thread team.
ordered	Required on a parallel for statement if an ordered directive is to be used in the loop.
schedule	Applies to the for directive.
nowait	Overrides the barrier implicit in a directive.

For data-sharing attributes:

Clause	Description
<code>private</code>	Specifies that each thread should have its own instance of a variable.
<code>firstprivate</code>	Specifies that each thread should have its own instance of a variable, and that the variable should be initialized with the value of the variable, because it exists before the parallel construct.
<code>lastprivate</code>	Specifies that the enclosing context's version of the variable is set equal to the private version of whichever thread executes the final iteration (for-loop construct) or last section (#pragma sections).
<code>shared</code>	Specifies that one or more variables should be shared among all threads.
<code>default</code>	Specifies the behavior of unscoped variables in a parallel region.
<code>reduction</code>	Specifies that one or more variables that are private to each thread are the subject of a reduction operation at the end of the parallel region.

Problem Statement 1: To Implement a parallel code for vector scalar addition

Screenshot 1:

```
Assignment-2 > C++ Que1_Addition_vector_scalar.cpp > main()
1  //Implement a parallel code for vector scalar addition
2  //2019BTECS00055 Aditya Manapure
3
4  #include<stdio.h>
5  #include<stdlib.h>
6  #include<omp.h>
7
8  #define THREAD_COUNT 8
9  #define ARRAY_SIZE 100
10 #define COUNT_PER_THREAD ARRAY_SIZE/THREAD_COUNT
11
12
13 int main(){
14
15     int A[ARRAY_SIZE], C[ARRAY_SIZE];
16     int scalar = 5;
17
18     for(int i=0; i<ARRAY_SIZE; i++){
19         A[i] = i;
20     }
21
22     omp_set_num_threads(THREAD_COUNT);
23     int i=0;
24     double start = omp_get_wtime();
25     #pragma omp parallel
26     {
27
28         #pragma omp single
29         printf("Number of threads Allocated %d\n", omp_get_num_threads());
30
31         #pragma omp single
32         printf("Processing Data ... ");
33
34         #pragma omp for
35         for(i=0; i<ARRAY_SIZE; i++){
36             printf("\n Thread id %d is processing on index %d\n", omp_get_thread_num(), i);
37             C[i] = A[i] + scalar;
38
39
40
41             #pragma omp barrier
42
43             #pragma omp single nowait
44             printf("Displaying Data ... ");
45
46             #pragma omp for
47             for(i=0; i<ARRAY_SIZE; i++){
48                 printf("\n Thread id %d is printing on index %d : %d + %d = %d", omp_get_thread_num(), i, A[i], scalar, C[i]);
49             }
50
51         }
52         double end = omp_get_wtime();
53         double total_time = end - start;
54         printf("\n\nTime taken : %lf\n", total_time);
55
56         return 0;
57     }
```

Output :

```

Thread id 2 is proccessing on index 34
Thread id 2 is proccessing on index 35
Thread id 2 is proccessing on index 36
Thread id 2 is proccessing on index 37
Thread id 2 is proccessing on index 38
displaying Data ...
Thread id 0 is printing on index 0 : 0 + 5 = 5
Thread id 0 is printing on index 1 : 1 + 5 = 6
Thread id 0 is printing on index 2 : 2 + 5 = 7
Thread id 0 is printing on index 3 : 3 + 5 = 8
Thread id 0 is printing on index 4 : 4 + 5 = 9
Thread id 0 is printing on index 5 : 5 + 5 = 10
Thread id 0 is printing on index 6 : 6 + 5 = 11
Thread id 0 is printing on index 7 : 7 + 5 = 12
Thread id 0 is printing on index 8 : 8 + 5 = 13
Thread id 0 is printing on index 9 : 9 + 5 = 14
Thread id 0 is printing on index 10 : 10 + 5 = 15
Thread id 0 is printing on index 11 : 11 + 5 = 16
Thread id 0 is printing on index 12 : 12 + 5 = 17
Thread id 7 is printing on index 88 : 88 + 5 = 93
Thread id 7 is printing on index 89 : 89 + 5 = 94

```

```

Thread id 5 is printing on index 67 : 67 + 5 = 72
Thread id 5 is printing on index 68 : 68 + 5 = 73
Thread id 5 is printing on index 69 : 69 + 5 = 74
Thread id 5 is printing on index 70 : 70 + 5 = 75
Thread id 5 is printing on index 71 : 71 + 5 = 76
Thread id 5 is printing on index 72 : 72 + 5 = 77
Thread id 5 is printing on index 73 : 73 + 5 = 78
Thread id 5 is printing on index 74 : 74 + 5 = 79
Thread id 5 is printing on index 75 : 75 + 5 = 80
Time taken : 0.094000

```

Information 1:

The clause `private (variable list)` indicates that the set of variables specified is local to each thread – i.e., each thread has its own copy of each variable in the list. The clause `firstprivate (variable list)` is similar to the `private` clause, except the values of variables on entering the threads are initialized to corresponding values before the `parallel` directive. The clause `shared (variable list)` indicates that all variables in the list are shared across all the threads, i.e., there is only one copy.

Problem Statement 2: To Implement a parallel code for vector-vector addition

Screenshot 1:

```
Assignment-2 > C++ Que2_Addition_vector_vector.cpp > main()
1 //Implement a parallel code for vector-vector addition
2 //Author : 2019BTECS00055 Aditya Manapure
3
4 #include<stdio.h>
5 #include<stdlib.h>
6 #include<omp.h>
7
8 #define THREAD_COUNT 8
9 #define ARRAY_SIZE 100
10 #define COUNT_PER_THREAD ARRAY_SIZE/THREAD_COUNT
11
12
13 int main(){
14
15     int A[ARRAY_SIZE], B[ARRAY_SIZE], C[ARRAY_SIZE];
16
17     for(int i=0; i<ARRAY_SIZE; i++){
18         A[i] = i;
19         B[i] = i;
20     }
21
22     omp_set_num_threads(THREAD_COUNT);
23     int i=0;
24     double start = omp_get_wtime();
25     #pragma omp parallel
26     {
27
28         #pragma omp single
29         printf("Number of threds Allocated %d\n", omp_get_num_threads());
30
31         #pragma omp single
32         printf("Processing Data ... ");
33
34         #pragma omp for
35         for(i=0; i<ARRAY_SIZE; i++){
36             printf("\n Thread id %d is proccessing on index %d\n", omp_get_thread_num(), i);
37             C[i] = A[i] + B[i];
38         }
39
40
41         #pragma omp barrier
42
43         #pragma omp single nowait
44         printf("Displaying Data ... ");
45
46         #pragma omp for
47         for(i=0; i<ARRAY_SIZE; i++){
48             printf("\n Thread id %d is printing on index %d : %d + %d = %d", omp_get_thread_num(), i, A[i], B[i], C[i]);
49         }
50     }
51     double end = omp_get_wtime();
52     double total_time = end - start;
53     printf("\n\nTime taken : %lf\n", total_time);
54
55     return 0;
56 }
57 }
```

Screenshot 2:

```
Thread id 5 is proccessing on index 70
Thread id 5 is proccessing on index 71
Thread id 5 is proccessing on index 72
Thread id 5 is proccessing on index 73
Thread id 5 is proccessing on index 74
Thread id 5 is proccessing on index 75
displaying Data ...
Thread id 5 is printing on index 64 : 64 + 64 = 128
Thread id 5 is printing on index 65 : 65 + 65 = 130
Thread id 5 is printing on index 66 : 66 + 66 = 132
Thread id 5 is printing on index 67 : 67 + 67 = 134
Thread id 5 is printing on index 68 : 68 + 68 = 136
Thread id 5 is printing on index 69 : 69 + 69 = 138
Thread id 5 is printing on index 70 : 70 + 70 = 140
Thread id 5 is printing on index 71 : 71 + 71 = 142
Thread id 5 is printing on index 72 : 72 + 72 = 144
Thread id 5 is printing on index 73 : 73 + 73 = 146
```

```
Thread id 0 is printing on index 6 : 6 + 6 = 12
Thread id 0 is printing on index 7 : 7 + 7 = 14
Thread id 0 is printing on index 8 : 8 + 8 = 16
Thread id 0 is printing on index 9 : 9 + 9 = 18
Thread id 0 is printing on index 10 : 10 + 10 = 20
Thread id 0 is printing on index 11 : 11 + 11 = 22
Thread id 0 is printing on index 12 : 12 + 12 = 24

Time taken : 0.093000
```

Information 2:

The schedule clause of the for directive deals with the assignment of iterations to threads. The general form of the schedule directive is `schedule (scheduling_class[, parameter])`. Open MP supports four scheduling classes: static, dynamic, guided, and runtime. The general form of the static scheduling class is `schedule (static[, chunk-size])`. This technique splits the iteration space into equal chunks of size `chunk-size` and assigns them to threads in a round-robin fashion. Open MP has a dynamic scheduling class. The general form of this class is `schedule (dynamic [, chunk-size])`. The iteration space is partitioned into chunks given by `chunk-size`. However, these are assigned to threads as they become idle.